

DIT407 Introduction to data science and AI

Assignment 6

Reynir Siik
reynir.siik@gmail.com

Franco Zambon
guszamfr@student.gu.se

2024-02-19

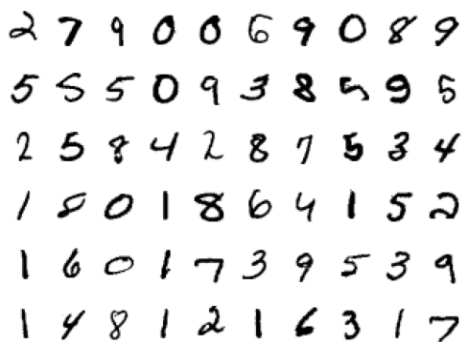
Abstract

This text is from the file `abstract.tex` and will be included in the abstract.

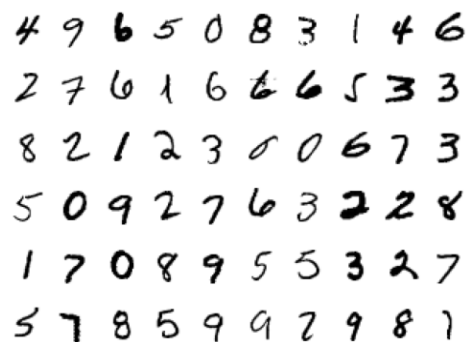
1 Problem 1:

The dataset is loaded with the specified tools, ref. [1].

The transforms function is used with a specifier for normalisation of the data in the range 0-1. The normalisation will use normal distribution with center 0.5 and $\sigma = 0.17$. This means that every (almost) value that falls within 3σ will fit into the specified range of 0-1. This loads a normalized dataset into Tensors as images wit 28 x 28 pixels. This was verified by printing the shape of the samples. See figure 1



(a) Images in the training set



(b) Images in the test set

Figure 1: Examples of the data set images

2 Problem 2: Single hidden layer

The neural network is initialised with the following parameters:

- input size = 784 (28X28 pixel of image)
- hidden layer = 284
- output = 10
- batch size = 2000
- learning rate = 0.01

When training the neural network the following table 3 shows the progress in learning.

Epoch	Accuracy
0	79.70
1	86.25
2	87.05
3	87.50
4	89.60
5	89.55
6	89.85
7	89.85
8	89.70
9	90.60
10	90.90
11	90.80
12	90.05
13	90.30
14	90.90
15	91.45

Table 1: Table showing accuracy after each epoch during training

3 Problem 3: Two hidden layers

3.1 Introduction

Essentially the same code is used for problem 3.

Line 21 is changed to 500

Line 22 is changed to 300

Line 93 is changed to 80

Line 78, 79, 89, 90 is uncommented.

For the optimizer, two parameters are introduced:

momentum = 0.9, and

weight decay = 0.001

See A.2

3.2 Results

The accuracy after each epoch is presented in table 2

To reach an accuracy over 98% it took 80 epochs of training.

Epoch	Accuracy	Epoch	Accuracy	Epoch	Accuracy	Epoch	Accuracy
0	87.15	20	96.45	40	97.50	60	98.55
1	90.20	21	95.80	41	97.05	61	97.90
2	91.35	22	96.50	42	97.20	62	98.00
3	91.40	23	97.00	43	97.80	63	97.25
4	91.90	24	96.95	44	97.35	64	98.15
5	92.15	25	96.80	45	97.70	65	97.85
6	92.90	26	96.75	46	97.35	66	97.65
7	94.05	27	97.00	47	97.75	67	97.95
8	92.95	28	96.70	48	97.40	68	97.95
9	93.90	29	96.95	49	97.20	69	97.95
10	93.90	30	97.30	50	97.45	70	98.15
11	94.55	31	97.85	51	98.10	71	97.60
12	94.45	32	97.05	52	98.20	72	97.95
13	95.15	33	97.35	53	97.50	73	98.30
14	95.10	34	96.65	54	97.45	74	98.05
15	96.20	35	96.75	55	97.25	75	98.35
16	95.10	36	96.80	56	97.75	76	97.60
17	95.50	37	97.20	57	97.55	77	98.00
18	96.50	38	97.35	58	97.65	78	98.10
19	96.20	39	96.90	59	97.95	79	98.45

Table 2: Training progress for 2 hidden layer neural network

Noticable is that the two hidden layers reach an accuracy of 96.2 after 16 epochs, whereas the single layer, which of course also is smaller, only had reached 91.45.

4 Problem 4: Convolutional neural network

4.1 Introduction

We implemented a simple convolutional neural network architecture for classifying MNIST digits. The network consisted of two convolutional layers followed by two fully connected layers, and it was trained using SGD with cross-entropy loss.

The training and validation loops were used to monitor the training progress and evaluate the model's performance on the test dataset.

4.2 The process and the results

First of all we defined the neural network architecture in the ConvNet class. We created 2 convolutional layers with a 3x3 kernel size and a padding of 1. We then created 2 fully connected layers, one with 128 neurons, the second with 10.

ReLU activation functions were applied after each convolutional layer and the first fully connected layer.

Max pooling layers with a kernel size of 2x2 and a stride of 2 are applied after each convolutional layer to reduce the spatial dimensions.

As before, we used ReLU, SGD, and cross entropy loss, and, to prevent overfitting, chose 0.01 as value for weight loss. We also chose a learning rate of 0.01. We then trained the

network for 40 epochs which seems to be enough since we reached our goal of 99% accuracy. The accuracy obtained is in table 3.

Epoch	Accuracy
0	92.82
1	96.04
2	97.20
3	97.61
4	97.58
5	98.28
6	98.38
7	98.34
8	98.45
9	98.65
10	98.71
11	98.70
12	98.66
13	98.77
14	98.76
15	98.80
16	98.77
17	98.81
18	98.93
19	98.92
20	98.90
21	98.91
22	98.96
23	98.91
24	99.01
25	98.99
26	99.07
27	99.02
28	99.05
29	98.98
30	99.08
31	99.06
32	98.95
33	99.08
34	99.01
35	99.06
36	99.02
37	98.98
38	98.97
39	99.00

Table 3: Table showing accuracy after each epoch during training

References

- [1] Y. LeCun, C. Cortes, and C. J. Burges, *The MNIST database of handwritten digits*, Retrieved 2024-01-02, 1998. [Online]. Available: <http://yann.lecun.com/exdb/mnist/>.

A Source code

This is the complete listing of the source code.

A.1 Problem 1 and 2

Code for splitting the datasets into training and test sets and training a linear model.

```
1 # Reynir Siik, Franco Zambon
2 # DIT407 lp3 2024-03-04
3 # Assignment 6
4 # Problem 1, 2
5 # Neural network
6 #
7 ## #####
8 import torch
9 from torch.utils.data import DataLoader
10 import torchvision
11 import matplotlib.pyplot as plt
12 import torch.nn as nn
13 from torchvision.transforms import transforms
14 ## #####
15 filelocation_in = r"Uppgift06\00_Uppg_Data\" # folder to get data from
16 filelocation_out = r"Uppgift06\test_data\" # folder to write data to
17 ## #####
18
19 ## initializing parameter
20 input_size=784 #28X28 pixel of image
21 hidden_size1=284 #size of 1st hidden layer(if used)
22 hidden_size=284 #size of last hidden layer
23 output =10 #output layer
24 batch_size=2000
25 lr_rate=0.01
26 #####
27 ## This transform should make the data normalized in the range 0-1 with
28 ## a probability of 0.97 (of the top of my head)
29 transform = transforms.Compose([transforms.ToTensor(),
30                                transforms.Normalize((0.5,), (0.17,)),
31                                ])
32
33 # Download test data from open datasets and normalize data in image
34 # Store the data locally so we don't have to download it everytime
35 # Training data
36 train_dataset = torchvision.datasets.MNIST(
37     root=filelocation_in,
38     train=True,
39     download=True,
40     transform=transform,
41 )
42
43 # Test data
44 test_data = torchvision.datasets.MNIST(
45     root=filelocation_in,
46     train=False,
47     download=True,
48     transform=transform,
49 )
50
```

```

51 # Create data loaders. Shuffle the data to avoid training on data in a
52 # certain order
53 train_dataloader = DataLoader(
54     train_dataset, batch_size=batch_size, shuffle=True
55 )
56 test_dataloader = DataLoader(
57     test_data, batch_size=batch_size, shuffle=True
58 )
59 #####
60 ## shape of the data that has been passed in batch
61 data=iter(train_dataloader)
62 samples,labels=next(data)
63
64 # Printing Images from dataset (quick and dirty)
65 # To be run twice with different setting to get both datasets
66 # figure = plt.figure()
67 # num_of_images = 60
68 # for index in range(1, num_of_images + 1):
69 #     plt.subplot(6, 10, index)
70 #     plt.axis('off')
71 #     plt.imshow(samples[index].numpy().squeeze(), cmap='gist_gray_r')
72 # plt.savefig(filelocation_out + 'A6_training_sample.pdf')
73
74 ## Set up the neural network, modified between problems
75 Mnist_model = nn.Sequential(
76     nn.Linear(input_size, hidden_size1),
77     nn.ReLU(),
78     # nn.Linear(hidden_size1, hidden_size),
79     # nn.ReLU(),
80     nn.Linear(hidden_size, output),
81     nn.LogSoftmax(dim=1),
82 )
83
84 loss=nn.CrossEntropyLoss()
85 # Recomendated parameters that have been found to work well after trying
86 # other values
87 optimizer=torch.optim.SGD(Mnist_model.parameters(),
88                             lr=lr_rate,
89                             # momentum=0.9,
90                             # weight_decay=0.05
91                             )
92
93 num_epochs = 16
94 print('Epoch┐,┐Accuracy')
95
96 for epoch in range(num_epochs):
97     for i,(images, labels) in enumerate(train_dataloader):
98         images = images.reshape(-1, 28*28)
99         ## forward connection
100         output = Mnist_model(images)
101         Loss = loss(output, labels)
102         ## calculating gradient
103         optimizer.zero_grad() # it will not cumulate the gradient result after every
104         Loss.backward()        # it will do backward propagation
105         optimizer.step()
106
107         # if (i+1)%batch_size == 0:
108         #     print(f"epoch={epoch+1}/{num_epochs}, step={i+1}/{len(train_dataloader)}")
109

```

```
110
111 ## model accuracy on validation dataset
112     with torch.no_grad():
113         n_correct=0
114         n_samples=0
115
116         for images,labels in test_dataloader:
117             images=images.reshape(-1,784)
118             output=Mnist_model(images)
119             labels=labels
120             _,prediction=torch.max(output,1)
121             n_samples=labels.shape[0]
122             n_correct+=(prediction==labels).sum().item()
123
124         accuracy=(n_correct/n_samples)*100
125
126     print(epoch, '\n', accuracy)
127     # print('Epoch: ',epoch, ', Accuracy = ', accuracy)
128
129 ## predict the result
130 predicted=[]
131 with torch.no_grad():
132     n_correct=0
133     n_samples=0
134     for images,labels in test_dataloader:
135         images=images.reshape(-1,784)
136         output=Mnist_model(images)
137         labels=labels
138         _,prediction=torch.max(output,1)
139         predicted.append(prediction)
140 print(prediction)
```

A.2 Problem 3

Code for two hidden layers.

```
1 # Reynir Siik, Franco Zambon
2 # DIT407 lp3 2024-03-04
3 # Assignment 6
4 # Problem 1, 2
5 # Neural network
6 #
7 ## #####
8 import torch
9 from torch.utils.data import DataLoader
10 import torchvision
11 import matplotlib.pyplot as plt
12 import torch.nn as nn
13 from torchvision.transforms import transforms
14 ## #####
15 filelocation_in = r"Uppgift06\00_Uppg_Data\\" # folder to get data from
16 filelocation_out = r"Uppgift06\test_data\\" # folder to write data to
17 ## #####
18
19 ## initializing parameter
20 input_size=784 # 28X28 pixel of image
21 hidden_size1=500 # size of 1st hidden layer(if used)
22 hidden_size=300 # size of last hidden layer
```



```
23 output =10          # output layer
24 batch_size=2000      # batch size
25 lr_rate=0.01         # learning rate
26 #####
27 ## This transform should make the data normalized in the range 0-1 with
28 ## a probability of 0.97 (of the top of my head)
29 transform = transforms.Compose([transforms.ToTensor(),
30                                transforms.Normalize((0.5,), (0.17,)),
31                                ])
32
33 # Download test data from open datasets and normalize data in image
34 # Store the data locally so we don't have to download it everytime
35 # Training data
36 train_dataset = torchvision.datasets.MNIST(
37     root=filelocation_in,
38     train=True,
39     download=True,
40     transform=transform,
41 )
42
43 # Test data
44 test_data = torchvision.datasets.MNIST(
45     root=filelocation_in,
46     train=False,
47     download=True,
48     transform=transform,
49 )
50
51 # Create data loaders. Shuffle the data to avoid training on data in a
52 # certain order
53 train_dataloader = DataLoader(
54     train_dataset, batch_size=batch_size, shuffle=True
55 )
56 test_dataloader = DataLoader(
57     test_data, batch_size=batch_size, shuffle=True
58 )
59 #####
60 ## shape of the data that has been passed in batch
61 data=iter(train_dataloader)
62 samples,labels=next(data)
63
64 # Printing Images from dataset (quick and dirty)
65 # To be run twice with different setting to get both datasets
66 # figure = plt.figure()
67 # num_of_images = 60
68 # for index in range(1, num_of_images + 1):
69 #     plt.subplot(6, 10, index)
70 #     plt.axis('off')
71 #     plt.imshow(samples[index].numpy().squeeze(), cmap='gist_gray_r')
72 # plt.savefig(filelocation_out + 'A6_training_sample.pdf')
73
74 ## Set up the neural network, modified between problems
75 Mnist_model = nn.Sequential(
76     nn.Linear(input_size, hidden_size1),
77     nn.ReLU(),
78     nn.Linear(hidden_size1, hidden_size),
79     nn.ReLU(),
80     nn.Linear(hidden_size, output),
81     nn.LogSoftmax(dim=1),
```

```
82 )
83
84 loss=nn.CrossEntropyLoss()
85 # Recommended parameters that have been found to work well after trying
86 # other values
87 optimizer=torch.optim.SGD(Mnist_model.parameters(),
88                             lr=lr_rate,
89                             momentum=0.9,
90                             weight_decay=0.001
91                             )
92
93 num_epochs = 80
94 print('Epoch□,□Accuracy')
95
96 for epoch in range(num_epochs):
97     for i,(images, labels) in enumerate(train_dataloader):
98         images = images.reshape(-1, 28*28)
99         ## forward connection
100        output = Mnist_model(images)
101        Loss = loss(output, labels)
102        ## calculating gradient
103        optimizer.zero_grad() # it will not cumulate the gradient result after every
104        Loss.backward()       # it will do backward propagation
105        optimizer.step()
106
107        # if (i+1)%batch_size == 0:
108        #     print(f"epoch={epoch+1}/{num_epochs}, step={i+1}/{len(train_dataloader)}")
109
110
111 ## model accuracy on validation dataset
112 with torch.no_grad():
113     n_correct=0
114     n_samples=0
115
116     for images,labels in test_dataloader:
117         images=images.reshape(-1,784)
118         output=Mnist_model(images)
119         labels=labels
120         _,prediction=torch.max(output,1)
121         n_samples=labels.shape[0]
122         n_correct=(prediction==labels).sum().item()
123
124     accuracy=(n_correct/n_samples)*100
125
126     print(f'{{epoch}}□,□{{accuracy:.2f}}')
127     # print('Epoch: ',epoch , ', Accuracy = ', accuracy)
128
129 ## predict the result
130 predicted=[]
131 with torch.no_grad():
132     n_correct=0
133     n_samples=0
134     for images,labels in test_dataloader:
135         images=images.reshape(-1,784)
136         output=Mnist_model(images)
137         labels=labels
138         _,prediction=torch.max(output,1)
139         predicted.append(prediction)
140 print(prediction)
```

A.3 Problem 4

Code for Non-linear relationship and multiple linear regression.

```
1 import torch.optim as optim
2 import torchvision.datasets as datasets
3
4 ## PROBLEM 4
5
6 ## Define the neural network architecture
7
8 class ConvNet(nn.Module):
9     def __init__(self):
10         super(ConvNet, self).__init__()
11         self.conv1 = nn.Conv2d(in_channels=1, out_channels=32, kernel_size=3, stride=1)
12         self.conv2 = nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3, stride=1)
13         self.fc1 = nn.Linear(64 * 7 * 7, 128)
14         self.fc2 = nn.Linear(128, 10)
15
16     def forward(self, x):
17         x = torch.relu(self.conv1(x))
18         x = torch.max_pool2d(x, kernel_size=2, stride=2)
19         x = torch.relu(self.conv2(x))
20         x = torch.max_pool2d(x, kernel_size=2, stride=2)
21         x = x.view(x.size(0), -1)
22         x = torch.relu(self.fc1(x))
23         x = self.fc2(x)
24         return x
25
26
27 ## Initialize the model and optimizer
28
29 model = ConvNet()
30 optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.9, weight_decay=0.01)
31
32
33
34 ## Training and validation
35
36 num_epochs = 40
37 for epoch in range(num_epochs):
38     model.train()
39     running_loss = 0.0
40     for batch_idx, (data, targets) in enumerate(train_loader):
41         optimizer.zero_grad()
42         outputs = model(data)
43         loss = criterion(outputs, targets)
44         loss.backward()
45         optimizer.step()
46         running_loss += loss.item()
47
48     print(f"Epoch [{epoch+1}/{num_epochs}], Loss: {running_loss/len(train_loader):.4f}")
49
50     model.eval()
51     correct = 0
52     total = 0
53     with torch.no_grad():
54         for data, targets in test_loader:
55             outputs = model(data)
```

```
56         _, predicted = torch.max(outputs.data, 1)
57         total += targets.size(0)
58         correct += (predicted == targets).sum().item()
59     print(f"Validation Accuracy: {(correct/total)*100:.2f}%")

1  # Reynir Siik, Franco Zambon
2  # DIT407 lp3 2024-02-11
3  # Assignment 4
4  # Problem 1, 2
5  # Life expectancy
6  # Non-linear relationship, Multiple linear regression
7  #####
```